HOV Image Detection

Ivan Chit Hin Ho (ivanho1@stanford.edu) - CS230 Winter 2019

Presentation: https://youtu.be/-PwwZRGBpV0
Source Code: https://github.com/alrightyi/hov

Motivation: Catch HOV-lane violators!!!



Image

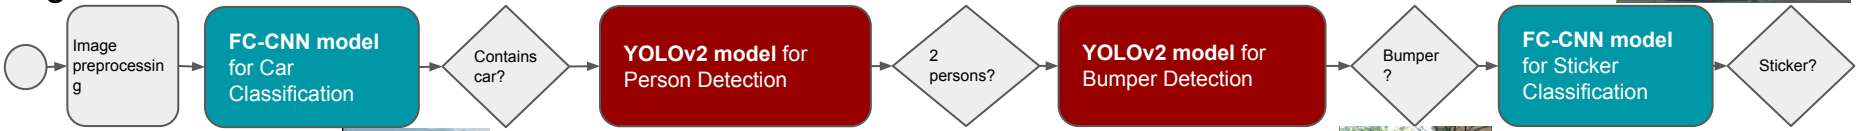

FC-CNN Model

Image size: 64x64x3
Batch size: 32
Epochs: 3
Optimizer: Adam
Loss: MSE
Train/Val: 80/20
Accuracy: 99.4%

(pre-trained model with dataset from ImageNet.)

Layer (type)	Output Shape	Param #
lambda_1 (Lambda)	(None, 64, 64, 3)	0
cv0 (Conv2D)	(None, 64, 64, 16)	448
dropout_1 (Dropout)	(None, 64, 64, 16)	0
cv1 (Conv2D)	(None, 64, 64, 32)	4640
dropout_2 (Dropout)	(None, 64, 64, 32)	0
cv2 (Conv2D)	(None, 64, 64, 64)	18496
max_pooling2d_1 (MaxPooling2	(None, 8, 8, 64)	0
dropout_3 (Dropout)	(None, 8, 8, 64)	0

pos images: 56, accuracy: 0.39285714285714285
neg images: 50, accuracy: 0.98
total images: 106, accuracy: 0.6698113207547169

YOLOv2 Model

Image size: 608x608x3
Batch size: 32
Epochs: 30
Optimizer: Adam
Loss: Classification & Coordinates Loss
Non-Max Suppression, IOU boxes

(pre-trained model with dataset from ImageNet.)

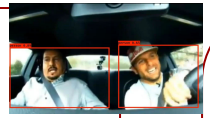

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 608, 608, 3)	0	
conv2d_1 (Conv2D)	(None, 608, 608, 32)	864	input_1[0][0]
batch_normalization_1 (BatchNor	(None, 608, 608, 32)	128	conv2d_1[0][0]
leaky_re_lu_1 (LeakyReLU)	(None, 608, 608, 32)	0	batch_normalization_1[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 304, 304, 32)	0	leaky_re_lu_1[0][0]
conv2d_2 (Conv2D)	(None, 304, 304, 64)	18432	max_pooling2d_1[0][0]
...	...	...	...
leaky_re_lu_21 (LeakyReLU)	(None, 38, 38, 64)	0	batch_normalization_21[0][0]
batch_normalization_20 (BatchNo	(None, 19, 19, 1024)	4096	conv2d_20[0][0]
space_to_depth_x2 (Lambda)	(None, 19, 19, 256)	0	leaky_re_lu_21[0][0]
leaky_re_lu_20 (LeakyReLU)	(None, 19, 19, 1024)	0	batch_normalization_20[0][0]
concatenate_1 (Concatenate)	(None, 19, 19, 1280)	0	space_to_depth_x2[0][0]

YOLOv2 Model

Image size: 416x416x3
Batch size: 32
Epochs: 30
Optimizer: Adam
Train/Val: 90/10
Loss: Classification & Coordinates Loss
Non-Max Suppression, IOU boxes, Early Stopping

Re-trained from YOLOv2 model using YAD2K + modifications.
Dataset from web search, hand-labeled using LabelBox (size: 680)

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 416, 416, 3)	0	
conv2d_1 (Conv2D)	(None, 416, 416, 32)	864	input_1[0][0]
batch_normalization_1 (BatchNor	(None, 416, 416, 32)	128	conv2d_1[0][0]
leaky_re_lu_1 (LeakyReLU)	(None, 416, 416, 32)	0	batch_normalization_1[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 208, 208, 32)	0	leaky_re_lu_1[0][0]
...	...	...	...
conv2d_22 (Conv2D)	(None, 13, 13, 1024)	11796480	concatenate_1[0][0]
batch_normalization_22 (BatchNo	(None, 13, 13, 1024)	4096	conv2d_22[0][0]
leaky_re_lu_22 (LeakyReLU)	(None, 13, 13, 1024)	0	batch_normalization_22[0][0]
conv2d_24 (Conv2D)	(None, 13, 13, 30)	30750	

FC-CNN Model

Image size: 64x64x3
Batch size: 32
Epochs: 3
Optimizer: Adam
Loss: MSE
Train/Val: 80/20
Accuracy: 67%
Dataset size: 1663 (1077 neg, 586 pos.)
Data Augment: flip, shift, rotate, shear, zoom, brightness

(pre-trained model with dataset from ImageNet.)

Layer (type)	Output Shape	Param #
lambda_1 (Lambda)	(None, 64, 64, 3)	0
cv0 (Conv2D)	(None, 64, 64, 16)	448
dropout_1 (Dropout)	(None, 64, 64, 16)	0
cv1 (Conv2D)	(None, 64, 64, 32)	4640
dropout_2 (Dropout)	(None, 64, 64, 32)	0
cv2 (Conv2D)	(None, 64, 64, 64)	18496
max_pooling2d_1 (MaxPooling2	(None, 8, 8, 64)	0
dropout_3 (Dropout)	(None, 8, 8, 64)	0